

```
In [1]: import datetime as dt
import os
import gc

import pandas as pd
# import dask.dataframe as dd
import numpy as np
import matplotlib.pyplot as plt

import lightgbm as lgb
from sklearn.model_selection import KFold
```

```
In [2]: pd.set_option('display.max_rows', 1000)
pd.set_option('display.max_columns', 100)
```

```
In [3]: INPUT_FOLDER = './data/input'
OUTPUT_FOLDER = './data/output'
```

```
In [4]: os.listdir(INPUT_FOLDER)
```

```
Out[4]: ['geo_params.csv',
'sample_final.csv',
'sku_final.csv',
'test_data.csv',
'train_final.csv']
```

Зчитано дані

```
In [5]: dateparser = lambda x: dt.datetime.strptime(x, '%Y-%m-%d')
```

```
In [6]: df_train = pd.read_csv(os.path.join(INPUT_FOLDER, 'train_final.csv'),
index_col='ID',
parse_dates = ['date'],
date_parser=dateparser)
```

```
In [7]: df_train['sales'].fillna(0, inplace=True)
```

```
In [8]: df_train.sort_values('date', inplace=True)
```

```
In [9]: df_train['price'] = df_train.groupby(['geoCluster', 'SKU'], sort=False)['price'].apply(
```

```
In [14]: df_test = pd.read_csv(
os.path.join(INPUT_FOLDER, 'test_data.csv'),
index_col='ID',
parse_dates = ['date'],
```

```
    date_parser=dateparser  
)
```

```
In [20]: df_train = df_train[df_train['date'] >= '2021-05-01']
```

```
In [21]: submission = pd.read_csv(os.path.join(INPUT_FOLDER, 'sample_final.csv'), index_col='ID'  
    submission.head()
```

```
Out[21]:
```

sales	
ID	
RR1666030	0
RR1666031	0
RR1666032	0
RR1666033	0
RR1666034	0

ID	
RR1666030	0
RR1666031	0
RR1666032	0
RR1666033	0
RR1666034	0

```
In [23]: geo_params = pd.read_csv(os.path.join(INPUT_FOLDER, 'geo_params.csv'))  
    geo_params.head()
```

```
Out[23]:
```

	geoCluster	cityId
0	21	1
1	47	1
2	48	1
3	92	1
4	112	1

	geoCluster	cityId
0	21	1
1	47	1
2	48	1
3	92	1
4	112	1

```
In [26]: sku = pd.read_csv(os.path.join(INPUT_FOLDER, 'sku_final.csv'))  
    sku.head()
```

```
Out[26]:
```

	SKU	productCategoryId	productCategory_caption_UKR	productCategory_caption_RU	productCategory
0	17	5416.0	Хурма	Хурма	
1	18	5413.0	Фейхоа	Фейхоа	
2	24	5425.0	Гранат	Гранат	
3	25	5431.0	Апельсин	Апельсин	

	SKU	productCategoryId	productCategory_caption_UKR	productCategory_caption_RU	productCategory
0	17	5416.0	Хурма	Хурма	
1	18	5413.0	Фейхоа	Фейхоа	
2	24	5425.0	Гранат	Гранат	
3	25	5431.0	Апельсин	Апельсин	

	SKU	productCategoryId	productCategory_caption_UKR	productCategory_caption_RU	productCategory_caption_ENG
4	208	5835.0	Вода України газована	Вода Украины газированная	Water, Uk

Data preprocessing

In [27]:

```
def merge_data(df):
    df = df.copy(deep=True)
    df = df.merge(sku, on = 'SKU', how='left')
    df = df.merge(geo_params, on = 'geoCluster', how='left')
    return df

def preprocess_date(df):
    df['month'] = df['date'].dt.month
    # df['weekofyear'] = df['date'].dt.weekofyear
    df['day'] = df['date'].dt.day
    df['dayofweek'] = df['date'].dt.dayofweek
    # df['dayofyear'] = df['date'].dt.dayofyear
    return df
```

In [28]:

```
df_train = merge_data(df_train)
```

In [29]:

```
df_train = preprocess_date(df_train)
```

In [30]:

```
df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 8514733 entries, 0 to 8514732
Data columns (total 28 columns):
#   Column                                Dtype
---  -
0   geoCluster                            int64
1   SKU                                  int64
2   date                                 datetime64[ns]
3   price                                float64
4   sales                                float64
5   productCategoryId                    float64
6   productCategory_caption_UKR          object
7   productCategory_caption_RU          object
8   productCategory_caption_ENG         object
9   productTypeId                        float64
10  productType_caption_UKR              object
11  productType_caption_RU              object
12  productType_caption_ENG             object
13  brandId                              float64
14  lagerUnitQuantity                    float64
15  lagerUnitTypeId                      int64
16  lagerUnitType_caption                object
17  trademark                            float64
18  countryOfOrigin                      float64
19  countryOfOrigin_caption              object
```

```

20 commodity_group          int64
21 commodity_group_caption_UKR object
22 commodity_group_caption_RU object
23 commodity_group_caption_ENG object
24 cityId                    int64
25 month                     int64
26 day                      int64
27 dayofweek                 int64
dtypes: datetime64[ns](1), float64(8), int64(8), object(11)
memory usage: 1.8+ GB

```

```
In [31]: # df_train.sample(15)
```

```
In [32]: cat_cols = [
    'geoCluster', 'SKU', 'productCategoryId', 'productTypeId', 'brandId', 'lagerUnitTyp
    'trademark', 'countryOfOrigin', 'commodity_group', 'cityId',
    'month',
    # 'weekofyear',
    'day',
    'dayofweek',
    # 'dayofyear'
]
num_cols = [
    'price',
    'lagerUnitQuantity',
    # 'price_change'
]
target = 'sales'
info_cols = ['date']
drop_cols = ['productCategory_caption_UKR', 'productCategory_caption_RU', 'productCateg
    'productType_caption_UKR', 'productType_caption_RU', 'productType_caption_ENG', 'la
    'commodity_group_caption_UKR', 'commodity_group_caption_RU', 'commodity_group_capti

def categorize_columns(df):
    df = df.copy(deep=True)
    df.drop(drop_cols, axis=1, inplace=True)
    df[cat_cols] = df[cat_cols].astype('category')
    return df

```

```
In [33]: df_train = categorize_columns(df_train)
```

```
In [34]: X, y = df_train.drop(target, axis=1).copy(), df_train[target].copy()
```

```
In [35]: del df_train
gc.collect()
```

```
Out[35]: 57
```

```
In [36]: sku_float_sales = X[y - y.astype(int) != 0]['SKU'].unique()
```

Model training

In [37]:

```
params = {
    'random_state':42,
    'n_estimators':500,
    'objective': 'regression',
    'metric': 'mae',
    "early_stopping_rounds":50,
    'learning_rate':0.1,
    'subsample':0.8,
    'subsample_freq':10,
    'feature_fraction':0.5
}
```

In [38]:

```
def custom_mae(y_true, y_pred, dates):
    df = pd.DataFrame(data={
        'y_true': np.asarray(y_true),
        'y_pred': np.asarray(y_pred),
        'date': np.asarray(dates)})
    df['error'] = np.abs(df['y_pred'] - df['y_true'])
    df_sum = df.groupby('date').sum().reset_index()
    df_sum.dropna(subset=['error', 'y_true'], how='any', inplace=True)
    df_sum = df_sum[df_sum['y_true'] > 0]
    if df_sum.shape[0] != df['date'].drop_duplicates().shape[0]:
        print(f"{df['date'].drop_duplicates().shape[0] - df_sum.shape[0]} rows were dro")
    return (df_sum['error'] / df_sum['y_true']).mean()
```

In [39]:

```
n_splits = 5
folds = np.array_split(X['date'].drop_duplicates().values, n_splits)
folds = [(fold[0], fold[-1]) for fold in folds]
for i, fold in enumerate(folds):
    print(f'Fold {i}: {fold}')
```

```
Fold 0: (numpy.datetime64('2021-05-01T00:00:00.000000000'), numpy.datetime64('2021-05-16
T00:00:00.000000000'))
Fold 1: (numpy.datetime64('2021-05-17T00:00:00.000000000'), numpy.datetime64('2021-06-01
T00:00:00.000000000'))
Fold 2: (numpy.datetime64('2021-06-02T00:00:00.000000000'), numpy.datetime64('2021-06-17
T00:00:00.000000000'))
Fold 3: (numpy.datetime64('2021-06-18T00:00:00.000000000'), numpy.datetime64('2021-07-03
T00:00:00.000000000'))
Fold 4: (numpy.datetime64('2021-07-04T00:00:00.000000000'), numpy.datetime64('2021-07-19
T00:00:00.000000000'))
```

In [40]:

```
models = []

train_mae_scores = []
val_mae_scores = []
```

In [41]:

```
for fold, fold_dates in enumerate(folds):

    print(f'\nFold {fold+1}')

    start_date, end_date = pd.Timestamp(fold_dates[0]), pd.Timestamp(fold_dates[1])
```

```

print(start_date, end_date)

train_mask = (X['date'] < start_date) | (X['date'] > end_date)
val_mask = (X['date'] >= start_date) & (X['date'] <= end_date)

X_train, y_train = X.loc[train_mask], y.loc[train_mask]
X_val, y_val = X.loc[val_mask], y.loc[val_mask]

train_dates, val_dates = X_train['date'], X_val['date']

X_train, X_val = X_train.drop('date', axis=1), X_val.drop('date', axis=1)

lgb_reg = lgb.LGBMRegressor(**params)
lgb_reg.fit(X_train, y_train, eval_set=[(X_train, y_train), (X_val, y_val)], verbose

y_train_pred = lgb_reg.predict(X_train)
y_val_pred = lgb_reg.predict(X_val)

#     y_train_pred = np.where(X_train['SKU'].isin(sku_float_sales), y_train_pred, np.ro
#     y_val_pred = np.where(X_val['SKU'].isin(sku_float_sales), y_val_pred, np.round(y_

train_mae_scores.append(custom_mae(y_train, np.round(y_train_pred), train_dates))
val_mae_scores.append(custom_mae(y_val, np.round(y_val_pred), val_dates))

print('Train mean daily MAE:', train_mae_scores[-1])
print('Valid mean daily MAE:', val_mae_scores[-1])

models.append(lgb_reg)

```

Fold 1

2021-05-01 00:00:00 2021-05-16 00:00:00

/home/eddie/miniconda3/envs/hack4retail/lib/python3.9/site-packages/lightgbm/sklearn.py:736: UserWarning: 'verbose' argument is deprecated and will be removed in a future release of LightGBM. Pass 'log_evaluation()' callback via 'callbacks' argument instead.

_log_warning("'verbose' argument is deprecated and will be removed in a future release of LightGBM. ")

[LightGBM] [Warning] early_stopping_round is set=50, early_stopping_rounds=50 will be ignored. Current value: early_stopping_round=50

[LightGBM] [Warning] feature_fraction is set=0.5, colsample_bytree=1.0 will be ignored. Current value: feature_fraction=0.5

/home/eddie/miniconda3/envs/hack4retail/lib/python3.9/site-packages/lightgbm/basic.py:1780: UserWarning: Overriding the parameters from Reference Dataset.

_log_warning('Overriding the parameters from Reference Dataset.')

/home/eddie/miniconda3/envs/hack4retail/lib/python3.9/site-packages/lightgbm/basic.py:1513: UserWarning: categorical_column in param dict is overridden.

_log_warning(f'{cat_alias} in param dict is overridden.')

[50] training's l1: 0.286592 valid_1's l1: 0.291598

Train mean daily MAE: 1.0370767630681577

Valid mean daily MAE: 1.0568761683609842

Fold 2

2021-05-17 00:00:00 2021-06-01 00:00:00

[LightGBM] [Warning] early_stopping_round is set=50, early_stopping_rounds=50 will be ignored. Current value: early_stopping_round=50

[LightGBM] [Warning] feature_fraction is set=0.5, colsample_bytree=1.0 will be ignored. Current value: feature_fraction=0.5

[50] training's l1: 0.283891 valid_1's l1: 0.299177

[100] training's l1: 0.281303 valid_1's l1: 0.299156

Train mean daily MAE: 1.0378611854219386

Valid mean daily MAE: 1.0314736089208234

Fold 3

2021-06-02 00:00:00 2021-06-17 00:00:00

[LightGBM] [Warning] early_stopping_round is set=50, early_stopping_rounds=50 will be ignored. Current value: early_stopping_round=50

[LightGBM] [Warning] feature_fraction is set=0.5, colsample_bytree=1.0 will be ignored. Current value: feature_fraction=0.5

```
[50] training's l1: 0.289658 valid_1's l1: 0.282616
[100] training's l1: 0.286938 valid_1's l1: 0.279262
[150] training's l1: 0.284991 valid_1's l1: 0.276907
[200] training's l1: 0.283635 valid_1's l1: 0.275909
[250] training's l1: 0.283071 valid_1's l1: 0.275302
[300] training's l1: 0.282165 valid_1's l1: 0.274934
[350] training's l1: 0.281606 valid_1's l1: 0.274417
[400] training's l1: 0.281006 valid_1's l1: 0.273919
[450] training's l1: 0.280188 valid_1's l1: 0.273133
[500] training's l1: 0.279418 valid_1's l1: 0.272655
```

Train mean daily MAE: 1.0211847411949384

Valid mean daily MAE: 1.045528178126395

Fold 4

2021-06-18 00:00:00 2021-07-03 00:00:00

[LightGBM] [Warning] early_stopping_round is set=50, early_stopping_rounds=50 will be ignored. Current value: early_stopping_round=50

[LightGBM] [Warning] feature_fraction is set=0.5, colsample_bytree=1.0 will be ignored. Current value: feature_fraction=0.5

```
[50] training's l1: 0.288168 valid_1's l1: 0.292
[100] training's l1: 0.285024 valid_1's l1: 0.291847
[150] training's l1: 0.283335 valid_1's l1: 0.291722
```

Train mean daily MAE: 1.0346770870699171

Valid mean daily MAE: 1.0828506211480438

Fold 5

2021-07-04 00:00:00 2021-07-19 00:00:00

[LightGBM] [Warning] early_stopping_round is set=50, early_stopping_rounds=50 will be ignored. Current value: early_stopping_round=50

[LightGBM] [Warning] feature_fraction is set=0.5, colsample_bytree=1.0 will be ignored. Current value: feature_fraction=0.5

```
[50] training's l1: 0.289072 valid_1's l1: 0.293897
[100] training's l1: 0.285566 valid_1's l1: 0.291556
[150] training's l1: 0.283666 valid_1's l1: 0.290518
[200] training's l1: 0.282289 valid_1's l1: 0.289072
[250] training's l1: 0.281734 valid_1's l1: 0.289713
```

Train mean daily MAE: 1.0311821878115128

Valid mean daily MAE: 1.097401858715899

```
In [42]: train_mae_scores, val_mae_scores = np.array(train_mae_scores), np.array(val_mae_scores)

print(f'Train score: {train_mae_scores.mean()}+-{train_mae_scores.std()}\n{train_mae_scores}')
print(f'Val score: {val_mae_scores.mean()}+-{val_mae_scores.std()}\n{val_mae_scores}\n')
```

Train score: 1.0323963929132929+-0.00606872006386897

[1.03707676 1.03786119 1.02118474 1.03467709 1.03118219]

Val score: 1.062826087054429+-0.024141512059155802

[1.05687617 1.03147361 1.04552818 1.08285062 1.09740186]

```
In [61]: pd.Series(index=X_train.columns, data=lgb_reg.feature_importances_).sort_values()
```

```
Out[61]: lagerUnitTypeId      5
commodity_group      7
countryOfOrigin      15
cityId      22
brandId      99
dayofweek      179
lagerUnitQuantity      185
trademark      220
month      290
productCategoryId      297
productTypeId      315
day      703
price      867
geoCluster      1371
SKU      1455
dtype: int32
```

```
In [45]: X_train.columns
```

```
Out[45]: Index(['geoCluster', 'SKU', 'price', 'productCategoryId', 'productTypeId',
               'brandId', 'lagerUnitQuantity', 'lagerUnitTypeId', 'trademark',
               'countryOfOrigin', 'commodity_group', 'cityId', 'month', 'day',
               'dayofweek'],
              dtype='object')
```

```
In [46]: gc.collect()
```

```
Out[46]: 14
```

Прогноз

```
In [47]: df_test = merge_data(df_test)
df_test = preprocess_date(df_test)
df_test = categorize_columns(df_test)
```

```
In [48]: X_sku_prices = X.groupby(['SKU'])['price'].agg(['min', 'max'])
```

```
In [49]: def adjust_price(price, sku):
    price_min_max = X_sku_prices.loc[sku]
    if price > price_min_max['max']:
        return price_min_max['max']
    if price < price_min_max['min']:
        return price_min_max['min']
    return price

train_skus = X_sku_prices.index.tolist()
df_test.loc[df_test['SKU'].isin(train_skus), 'price'] = df_test.loc[df_test['SKU'].isin
```

```
In [50]: df_test.drop(info_cols, axis=1, inplace=True)
```


```
In [51]:
```



```
predictions = []  
for model in models:  
    predictions.append(model.predict(df_test, num_iteration=lgb_reg.best_iteration_))
```


In [52]: `y.hist()`

Out[52]: <AxesSubplot:>



In [53]: `submission['sales'] = np.array(predictions).mean(axis=0)`
`submission['sales'].hist()`

Out[53]: <AxesSubplot:>




In [54]: `submission['sales'].min()`

Out[54]: -2.7819200747371204

In [55]: `submission.loc[submission['sales']<0, 'sales'] = 0`
`submission['sales'] = np.round(submission['sales'])`
`submission['sales'].hist()`

Out[55]: <AxesSubplot:>



In [57]:

```
ts = dt.datetime.now().strftime('%Y%m%d_%H_%M_%S')
submission.to_csv(
    os.path.join(
        OUTPUT_FOLDER,
        f'{ts}.csv'
    )
)
```